

[Home](#)[Intro](#)[Docs](#)[Github](#)

# Getting started

MapDB has very power-full API, but for 99% cases you need just two classes: [DBMaker](#) is builder style factory for configuring and opening a database. It has handful of static 'newXXX' methods for particular storage mode. [DB](#) represents storage. It has methods for accessing Maps and other collections. It also controls DB life-cycle with commit, rollback and close methods.

Best place to checkout various features of MapDB are [Examples](#). There is also [screencast](#) which describes most aspects of MapDB.

There is [MapDB Cheat Sheet](#), on just two pages it is quick reminder of MapDB capabilities.

## Maven

MapDB is in Maven Central. Just add code bellow to your pom file to use it. You may also download jar file directly from [repo](#).

```
<dependency>
  <groupId>org.mapdb</groupId>
  <artifactId>mapdb</artifactId>
  <version>1.0.7</version>
</dependency>
```

[ChangeFileRegex](#)

We are working on new generation of MapDB. It is faster and more reliable. Latest semi-stable build is at [snapshot repository](#):

```
<repositories>
  <repository>
    <id>sonatype-snapshots</id>
    <url>https://oss.sonatype.org/content/repositories/snapshots</url>
  </repository>
</repositories>

<dependencies>
  <dependency>
    <groupId>org.mapdb</groupId>
    <artifactId>mapdb</artifactId>
    <version>2.0.0-SNAPSHOT</version>
  </dependency>
</dependencies>
```

[ChangeFileRegex](#)[ChangeFileRegex](#)

## Hello World

Hereafter is a simple example. It opens TreeMap backed by file in temp directory, file is discarded after JVM exit:

```
import org.mapdb.*;
ConcurrentNavigableMap treeMap = DBMaker.newTempTreeMap()

// and now use disk based Map as any other Map
treeMap.put(111, "some value")
```

[RootFile](#)

More advanced example with configuration and write-ahead-log transaction.

```
import org.mapdb.*;

// configure and open database using builder pattern.
// all options are available with code auto-completion.
DB db = DBMaker.newFileDB(new File("testdb"))
    .closeOnJvmShutdown()
```

```
.encryptionEnable("password")
.make();

// open existing an collection (or create new)
ConcurrentNavigableMap<Integer,String> map = db.getTreeMap("collectionName");

map.put(1, "one");
map.put(2, "two");
// map.keySet() is now [1,2]

db.commit(); //persist changes into disk

map.put(3, "three");
// map.keySet() is now [1,2,3]
db.rollback(); //revert recent changes
// map.keySet() is now [1,2]

db.close();
```

**RootFile**

## What you should know

MapDB is very simple to use, however it bites when used wrong way. Here is list of most common usage errors and things to avoid:

- Transactions (write-ahead-log) can be disabled with `DBMaker.transactionDisable()`, this will speedup writes. However without transactions store gets corrupted when not closed correctly.
- Keys and values must be immutable. MapDB may serialize them on background thread, put them into instance cache... Modifying an object after it was stored is a bad idea.
- MapDB relies on memory mapped files. On 32bit JVM you will need `DBMaker.randomAccessFileEnable()` configuration option to access files larger than 2GB. RAF introduces overhead compared to memory mapped files.
- MapDB does not run defrag on background. You need to call `DB.compact()` from time to time.
- MapDB uses unchecked exceptions. All `IOException` are wrapped into unchecked `IOError`. MapDB has weak error handling and assumes disk failure can not be recovered at runtime. However this does not affects data safety, if you use durable commits.